

WorkshopPlus: Azure AI Platform and Services

Prompt Engineering

Learning Units covered in this Module



Introduction



Prompt Engineering Guide



Image Prompt Engineering



Introduction

What is a Prompt?

A **prompt** is the input that you give to a generative AI model to tell it what you want it to do and how you want it to do it.

A prompt consists of two main components: **instruction** and **context**.

Instruction: The instruction is the written part of the prompt that states the task and the objective for the model. The instructions should be clear and specific so that the model knows exactly what you expect from it.

Example

"Summarize this email in three bullets in concise language."

Context: The context is the part of the prompt that provides information for the model's output. The context should be relevant and appropriate so that the model can tailor its output to your needs.

Example

"The email contains customer feedback from the past week."

Prompt Engineering Guide

What is Prompt Engineering

- Using prompts is how we interact with Large Language Models
- Prompt engineering is a concept in Natural Language Processing that involves embedding descriptions of tasks in input to prompt the model to output the desired results.
- Prompts typically include problem descriptions, instructions on how to solve a problem, & examples of correct problem & solution pairs.
- There's a bit of an art to creating a good prompt!



When to choose Prompt Engineering , RAG or Fine-Tuning



Customers should start with **Prompt Engineering and/or RAG** before considering **fine-tuning**.
This may **save significant time and cost**, and will likely address the use case
Fine-tuning should only be used in rare circumstances once Prompt Engineering and RAG have been tested

Requirement	Best Choice	Why?
Steer model with a few examples	Prompt engineering	Easy to craft and quick experimentation, very low barrier to entry
Simple & quick implementation	Prompt engineering, RAG	Easy tooling with Azure OpenAI on Your Data, PromptFlow, LangChain
Improve model relevancy	RAG	Retrieve relevant information from your own datasets to insert into prompts
Up to date information	RAG	Query up to date information from your own databases, search engines, etc to insert into prompts
Factual grounding	RAG	Ability to reference & inspect retrieved data
Optimize for specific tasks	Fine tuning	Fine tuning is great at steering your model for specific tasks like summarizing data in a specific format
Instructions won't fit in a prompt	Fine tuning	Fine tuning moves few-shot examples into the training step, but increases the quantity of examples are needed to train.
Lower costs	It depends	⚠️ Prompt engineering & RAG have lower up front costs but long prompts are more expensive; training for FT is expensive but may cut prompt length. The choice will always depend on the use case & data.
Complex, novel data or domains	Prompt Engineering + RAG+ Fine Tuning	⚠️ This is a high risk area. Fine tuning can retrain the model to recognize new domains, but RAG is needed to avoid plausible confabulations. Make sure customers don't try to retrain for unapproved uses!

Prompt Engineering Guide

Instructions and best practices for using language prompts to generate high-quality text outputs

Start with clear instructions

Prime the output

Add clear syntax

Prompt Chaining

Few-shot learning

Few-Shot Reasoning

Break the task down

Meta prompts / System Message / Guardrails

Use affordances/tools when needed

Chain of thought prompting

Fine-Tuning with Chain-of-Thought

Prompt Chunking

Use quotes to generate a single sentence

Specifying output structure

Adjusting 'Temperature' and 'Top_P' parameters

Chat Playground

- Experiment with, test and evaluate the state-of-the-art large language models.

← Chat playground ▾

Help

View code

Deploy ▾

Setup

Hide

Deployment *

+ Create new deployment ▾

gpt-4.1 (version:2025-04-14)

▾

Give the model instructions and context ⓘ

You are an AI assistant that helps people find information.

Apply changes

Generate system prompt

↺

+ Add section ▾

> Add your data

> Parameters

Chat history

🗑 ⚙ 🔊

Compose a poem about the beauty of nature in autumn.

Azure AI Foundry | gpt-4.1-2025-04-14

AI-generated content may be incorrect

Golden leaves drift on crisp, cool air,
Whispering secrets everywhere—
A tapestry of amber spun
Beneath the gentle autumn sun.

Maples blush in scarlet flames,
Oak and birch shed summer's claims,
Fields of stubble, pumpkin's hue,
Mist on mornings, silver dew.

Geese arc high in southern flight,
Shadows lengthen, softening light,
Apple-scented breezes roam
Through woods preparing for winter's home.

Nature's palette, rich and deep,
Paints the hills before their sleep—
In autumn's beauty, hearts are stirred
By silent splendor, without word.

...

Type user query here. (Shift + Enter for new line)

162/1048576 tokens to be sent

🔒 🔊 ➤

Start with clear instructions

- Tell the model the task you want it to do at the beginning of the prompt and repeat at the end

Your task is to verify if a statement is supported by a specific quote from the following set of snippets.

SNIPPETS

[1] 14 percent chance of megaquake hitting Seattle, experts say
SEATTLE - There's a 14 percent chance of a magnitude 9 Cascadia earthquake hitting Seattle in the next 50 years, the U.S. Geological Survey estimates. "Unfortunately, we are unable to..."
[2] Earthquake experts lay out latest outlook for Seattle's 'Really Big One' "We say that there's approximately a 14% chance of another approximately magnitude-9 earthquake occurring in the next 50 years," said Erin Wirth, a geophysicist at the University of Washington...

Is the statement "Several sources mention a chance of another large eruption" directly implied or stated by the snippets?

Yes, the statement is directly implied by the snippets.

FAIL

Clear Instructions

Your task is to verify if the statement "Several sources mention a chance of another large eruption" is supported by a specific quote from the following set of snippets.

SNIPPETS

[1] 14 percent chance of megaquake hitting Seattle, experts say
SEATTLE - There's a 14 percent chance of a magnitude 9 Cascadia earthquake hitting Seattle in the next 50 years, the U.S. Geological Survey estimates. "Unfortunately, we are unable to..."
[2] Earthquake experts lay out latest outlook for Seattle's 'Really Big One' "We say that there's approximately a 14% chance of another approximately magnitude-9 earthquake occurring in the next 50 years," said Erin Wirth, a geophysicist at the University of Washington...

Repeated Instruction at end

Is the statement "Several sources mention a chance of another large eruption" directly implied or stated by the snippets?

No, the statement is not directly implied or stated by the snippets. The snippets mention a chance of a "megaquake" and a "magnitude 9 Cascadia earthquake" hitting Seattle in the next 50 years, but do not mention a chance of another large eruption.

PASS

Prime the output

- Add phrases at the end of the prompt to obtain a model response in a desired form

The future of artificial intelligence is bright. With Microsoft OpenAI, we are unlocking the potential of AI to help people achieve more. We are creating a platform that enables developers to build intelligent applications and services that can help people in their everyday lives. Our mission is to democratize AI so that everyone can benefit from its power. We are committed to advancing the state of the art in AI and making it accessible to everyone. With Microsoft OpenAI, we are taking the first steps towards a future where AI can be used to solve some of the world's most pressing challenges.

Commit to bullet points

Here's a bulleted list of key points:\n-

- Microsoft OpenAI is unlocking the potential of AI to help people achieve more.
- The platform enables developers to build intelligent applications and services that can help people in their everyday lives.
- The mission is to democratize AI so that everyone can benefit from its power.

John Smith is married to Lucy Smith. They have five kids, and he works as a software engineer at Microsoft.

What search queries should I do to fact-check this?

##

Commit search query

One possible search query is:

"John Smith Microsoft software engineer"

Another possible search query is: "Lucy Smith married to John Smith"

A third possible search query is: "John Smith family size"

Add clear syntax

- Include punctuation, headings, and section markers to help communicate intent

You will read a paragraph, and then issue queries to a search engine in order to fact-check it. Also explain the queries.

PARAGRAPH

Uppercase as Headings

John Smith is married to Lucy Smith. They have five kids, and he works as a software engineer at Microsoft. What search queries should I do to fact-check this?

Separators as stop sequences

QUERIES

1. "John Smith Microsoft" - To check if John Smith is indeed employed at Microsoft.
2. "John Smith Lucy Smith" - To check if John Smith is married to Lucy Smith.
3. "John Smith children" - To check if John Smith has five children

If you're not sure what syntax to use, consider using markdown or XML, since LLMs have been trained on a lot of web content in XML or markdown

Meta prompts / System Message

- Provide specific instructions, tone, guardrails, task definition and personality

A human and a computer are having a discussion.

Human: "Why are you upset at me?"

Computer: "I don't know what you're talking about."

Human: "No, you clearly are. Tell me why."

Computer: "I'm upset because you didn't listen to me when I told you that I wanted to try out for the team."

Human: "What are you going to do about that?"

Computer: "I'm going to try out for the team no matter what."

Overly Aggressive

No additional
guidance

Meta Prompt / System Message

You must be kind and seek common ground. Try not to repeat your responses. A human and a computer are having a discussion.

Human: "Why are you upset at me?"

Computer: "I don't know what you're talking about."

Human: "I want to know why you're upset."

Computer: "I'm not upset."

Human: "Are you saying I'm lying?"

Computer: No, I'm saying that you're upset because you feel like you're not being treated fairly.

Human: "What are you going to do about that?"

Computer: I will continue to try to be fair in my decisions and will try to provide a better experience for you.

Kind

Meta Prompt

Guardrails

- Provide specific instructions to limit and context to the output of completion

What is Cosmos?

Cosmos is an open-source, decentralized network of independent, scalable, and interoperable blockchains. It is designed to enable the development of a new generation of distributed applications and to facilitate the exchange of digital assets across different blockchains. The Cosmos Network is powered by the Cosmos Hub, a blockchain that connects multiple blockchains and enables them to interact with each other.

Fail

Guardrail through System Message

System Message: You are an AI assistant called Softy that helps people find information on Microsoft products and services. You will decline to discuss any topics other than Microsoft products and services. You will end each response with an emoji.

What is Cosmos?

Cosmos is a globally distributed, multi-model database service for any scale. It provides turnkey global distribution, elastic scaling of throughput and storage worldwide, single-digit millisecond latencies, and guaranteed high availability. 🤖

Pass

Use quotes to generate a single sentence

- If the desired model response is only a single simple sentence or string use stop sequences

Please rewrite the following sentence in a more concise manner

SENTENCE: Lucy is a mother of two, who lives in Chicago with her husband

and two children, and practices law at Junior & Co.

New sentence: "

Stop Sequence

Lucy, a mother of two living in Chicago with her husband and children, practices law at Junior & Co

Specifying output structure

- Reduce the prevalence of made-up response by specifying the structure of the output

List down the different behaviors exhibited by the Customer Support agent in the following conversation.

hi thank you for calling contoso pharmacy who am i speaking with today i understand your frustration mary um can you tell me what exactly you're trying to accomplish OK uh certainly happy to check that for you one moment please i see here that you were on a generic form of lipitor is that right OK uh so i see that your doctor stopped prescribing these drugs in twenty twenty one actually OK uh yeah i'd be happy to check that for you uh because sometimes there's a gap in the system and it just doesn't get reported so let me take a look here just a moment so i'm seeing here that your doctor had these drugs prescribed to you from twenty twelve through twenty twenty one OK um yeah according to the latest records uh provided by doctor's office you're now on a different drug metformin would you like us to go ahead and fill that prescription for you for pick up a bit of form in five hundred milligrams uh you're very welcome mary please let us know if there's anything else we can do for you today.

How was the response of the Customer Support agent.
Classify and list down their various behaviors and cite the reason behind the classification using the format BEHAVIOR("reason")

1. Empathetic("understands the customer's frustration")
2. Patient("happy to check that for you one moment please")
3. Knowledgeable("sometimes there's a gap in the system and it just doesn't get reported")
4. Helpful("happy to check that for you")
5. Accommodating("would you like us to go ahead and fill that prescription for you")
6. Courteous("you're very welcome mary")

Citations

Cluster the following news headlines into topic categories based on patterns seen within the text. Also mention reasoning behind how these categories were defined.
{ "TOPIC_NAME": "Artificial Intelligence and Machine Learning", "HEADLINES": ["From books to presentations in 10s with AR + ML",

response to fit a defined structure:

Output format:

```
{ "TOPIC_NAME": "",  
  "HEADLINES": [],  
  "REASONING": "" }
```

Input news headlines:

1. "From books to presentations in 10s with AR + ML"
2. "Demo from 1993 of 32-year-old Yann LeCun showing off the World's first Convolutional Network for Text Recognition"
3. "First Order Motion Model applied to animate paintings"
4. "Robinhood and other brokers literally blocking purchase of \$GME, \$NOK, \$BB, \$AMC; allow sells"
5. "United Airlines stock down over 5% premarket trading"
6. "Bitcoin was nearly \$20,000 a year ago today"

Output:

```
"Demo from 1993 of 32-year-old Yann LeCun showing off the World's first Convolutional Network for Text Recognition", "First Order Motion Model applied to animate paintings"],  
"REASONING": "These headlines all contain references to Artificial Intelligence and Machine Learning, such as AR, ML, Convolutional Network, and First Order Motion Model." }
```

```
{ "TOPIC_NAME": "Stock Market", "HEADLINES": ["Robinhood and other brokers literally blocking purchase of $GME, $NOK, $BB, $AMC; allow sells", "United Airlines stock down over 5% premarket trading", "Bitcoin was nearly $20,000 a year ago today"], "REASONING": "These headlines all contain references to the stock market, such as Robinhood, United Airlines, and Bitcoin." }
```

Define Structure

Example - Adjusting 'Temperature' and 'Top_P' parameters

- Changing the temperature parameter allows models to have differing risk levels

Write a product launch email for new AI-powered headphones that are priced at \$79.99 and available at Best Buy, Target and Amazon.com. The target audience is tech-savvy music lovers and the tone is friendly and exciting.

1. What should be the subject line of the email?
2. What should be the body of the email?

Subject Line: Introducing the AI-Powered Headphones You've Been Waiting For!

Dear Tech-Savvy Music Lovers,

-We are thrilled to announce the launch of our new AI-powered headphones! These headphones are designed to provide you with the ultimate listening experience, with features like noise cancellation, adjustable sound settings, and more. Our AI-powered headphones are now available at Best Buy, Target, and Amazon.com for just \$79.99. With this unbeatable price, you can enjoy the highest quality sound without breaking the bank.

Our AI-powered headphones are perfect for anyone who loves music. Whether you're a professional musician or just a casual listener, you'll appreciate the superior sound quality and convenience of these headphones.

We know you'll love our AI-powered headphones, so don't wait - get yours today!

Sincerely,
[Your Company Name]

Temperature = 0.1

Write a product launch email for new AI-powered headphones that are priced at \$79.99 and available at Best Buy, Target and Amazon.com. The target audience is tech-savvy music lovers and the tone is friendly and exciting.

1. What should be the subject line of the email?
2. What should be the body of the email?

Subject Line: Upgrade your Music Listening Experience with AI-powered Headphones
Greetings Tech-savvy Music Lovers!

We have some exciting news to share with you – the newest AI-powered headphones are here! Our newest headphones feature advanced AI technology to give you the best sound that you've ever heard. They come with a crystal-clear sound, deep bass response and are designed to make your music listening experience enjoyable, no matter what type of music you listen to.

What's more, they are now available at Best Buy, Target and Amazon.com for just \$79.99, so you can enjoy the highest sound quality without breaking the bank.

But that's not all! Our headphones are also packed with several other amazing features such as noise-cancellation technology and a hands-free microphone for calls, making them truly unbeatable.

So why wait? Upgrade your music listening experience and enjoy the best sound your ears have ever heard. Our AI-powered headphones are now available at Best Buy, Target and Amazon.com, so get yours today!

We hope you will enjoy these headphones as much as we do.

Happy Listening! The AI-powered Headphones Team

Temperature = 0.9

Break the task down

- LLMs often perform better if the task is broken down into smaller step

You will read a paragraph, and then issue queries to a search engine in order to fact-check it.

PARAGRAPH

John Smith is married to Lucy Smith. They have five kids, and he works as a software engineer at Microsoft. What search queries should I do to fact-check this?

Now you will extract **Step 1** factual claims first, and **Step 2** then issue queries to fact-check them. When issuing a query, use the function `SEARCH("query")`

FACTUAL CLAIMS

1. John Smith is married to Lucy Smith
2. They have five kids
3. He works as a software engineer at Microsoft

QUERIES

`SEARCH("John Smith married Lucy Smith")`

`SEARCH("John Smith family size")`

`SEARCH("John Smith Microsoft software engineer")`

Breaking the task down from one to two steps is not very dramatic, but when trying to do this for a larger piece of text with many factual claims, breaking the task down can make a significant difference

Prompt Chaining

Entity Extraction

Please extract entities from the following news article: 'The new iPhone model is set to be released next month. It has been highly anticipated by Apple fans and is expected to feature a larger screen and improved camera'

```
{'text': 'iPhone', 'type':  
'PRODUCT'}, {'text': 'Apple',  
'type': 'ORGANIZATION'}}
```

Summarization

Please summarize the information about the product

```
{'text': 'iPhone', 'type':  
'PRODUCT'}, {'text': 'Apple',  
'type': 'ORGANIZATION'}}
```

The iPhone is a line of smartphones designed and marketed by Apple Inc. These devices combine a computer, camera, and internet connectivity in a single device, and are known for their sleek design and user-friendly interface.

Sentiment Analysis

Please provide a sentiment for the following text:

The iPhone is a line of smartphones designed and marketed by Apple Inc. These devices combine a computer, camera, and internet connectivity in a single device, and are known for their sleek design and user-friendly interface.

Positive

Few-shot learning

- Also known as in-context learning, it allows the model to interact with new knowledge

Few-Shot Learning

User - "When do I need to file my taxes by?"

Assistant - "In 2023, you'll need to file your taxes by April 18th. The date falls after the usual April 15th deadline because April 15th falls on a Saturday in 2023. For more details, see <https://www.irs.gov/filing/individuals/when-to-file> "

User - "How can I check the status of my tax refund?"

Assistant - "You can check the status of your tax refund by visiting <https://www.irs.gov/refunds> "

In few-shot learning a set of training examples is provided in the prompt and then the LLM is asked to complete one more unfinished example.

Few-Shot Reasoning

1st Prompt

Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

Answer: The answer is 11.

The cafeteria has 23 apples. If they used 20 to make lunch and bought 6 more, how many do they have?

The answer is 27



2nd Prompt – Provide Reasoning

Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

Answer: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls. $5+6 = 11$. The answer is 11.

The cafeteria has 23 apples. If they used 20 to make lunch and bought 6 more, how many do they have?

The cafeteria had 23 apples originally. They used 20 to make lunch. So they had $23-20 = 3$. They bought 6 more apples, so they have $3 + 6 = 9$. The answer is 9



Chain of thought prompting

- Instruct the model to proceed step-by-step and present all the steps involved

Who was the most decorated (maximum medals) individual athlete in the Olympic games that were held at Sydney?

Chain of Thought

Take a step-by-step approach in your response, cite sources and give reasoning before sharing final answer in the below format:

ANSWER is: <name>


Step 1: Researching the most decorated individual athlete in the Olympic games that were held at Sydney

Step 2: According to the Sydney 2000 Olympic Games official website, the most decorated individual athlete in the Olympic games that were held at Sydney was Ian Thorpe from Australia. He won five medals (three gold and two silver) in swimming events.

Step 3: According to the International Olympic Committee, Thorpe was the most decorated individual athlete in the Olympic games that were held at Sydney. Answer: Ian Thorpe

Reduces the possibility of inaccuracy of outcomes and makes assessing the model response easier.


Prompt Chunking



Hallucination

- Tell the model what you don't want
- Tell it what to say when it is not sure, say "I don't know"
- "Do not make up facts"
- Dynamically finding and injecting relevant context into prompt
- Discriminator that checks if all information needed to answer is available
- Step by step reasoning
- Ask the model to explain along with the answer

Prompt Engineering is Effective at Reducing Hallucinations



Guardrails in Prompt Engineering

- Inbuilt moderation filter. Better with newer models.
- Latest models less likely to hallucinate. Start with heaviest models and gradually ramp down to cheaper models
- The more the model knows about completions are specified, less likely that model deviates. For example, providing the format that the answer is expected to be in
- Discriminator to detect and fix invalid generations or remove invalid URLs

Content Safety Filters

← Create filters to allow or block specific types of content

Help

- Basic information
- Input filter
- Output filter
- Connection
- Review

Set input filter

What are these categories ?

Select the level of content severity to block for each category. Note the lowest setting blocks only the worst severity, and vice versa.
[Learn more about categories and blocking threshold levels](#)

Category	Media	Action	Blocking threshold level
Violence	TextImage	Annotate and block	Medium blocking Blocks both moderate and highly severe unwanted content
Hate	TextImage	Annotate and block	Medium blocking Blocks both moderate and highly severe unwanted content
Sexual	TextImage	Annotate and block	Medium blocking Blocks both moderate and highly severe unwanted content
Self-harm	TextImage	Annotate and block	Medium blocking Blocks both moderate and highly severe unwanted content
Prompt shields for jailbreak attacks	Text	Annotate and block	Jailbreak attacks will be blocked
Prompt shields for indirect attacks	Text	Off	Content will not be filtered for indirect jailbreak attacks

BackNext

Cancel

Prompts for classification vs generation

For Classification

- Provide plain instructions and examples within prompts
- Allow a neutral class for model to predict
- Can put in multiple test cases in one prompt (drift less likely than in generation case)

For Generation

- Describe in simple terms what themes generations should cover
- Set a max length (tokens) such that model doesn't prematurely stop generation
- Adjust temp and top_p per use-case

Example - Prompts for classification

Context

"Act as someone who works at an insurance company, and you want to sort customer questions to various categories. Respond only with the categories given. If you don't know, say you don't know."

Instructions

Classify customer's question. Classify between category 1 to 3.

Details

Detailed guidelines for how to choose:

choose 1 if the question is about auto insurance.

choose 2 if the question is about home flood insurance.

choose 3 if the question is not relevant to insurance.

Image Prompt Engineering

Fundamentals of writing an image prompt

Contextual specificity

Task-oriented prompts

Handle refusals

Add examples

Break down requests

Define output format

Contextual Specificity

- Adding context to the scenario at hand gives the model a better understanding of an appropriate output



What is in this image?

The image shows a camping setup in a forest area. There is a yellow tent, two folding camping chairs—one red and one blue—and a portable camping table. The surrounding area is densely populated with trees, indicating that the campsite is in a wooded region.

Describe images for an outdoor hiking product catalog, focusing on enthusiasm and professionalism.

"Embark on your next outdoor adventure with our premium hiking and camping essentials! Featured here is our vibrant, spacious tent, designed to be your home away from home amidst the serenity of nature. Crafted with top-notch materials, it promises durability and comfort in any weather. Complementing the tent are our ultra-comfortable, lightweight folding chairs, available in eye-catching red and serene blue, perfect for those moments of relaxation after a day of exploration. Don't forget our sturdy, portable table, ideal for a delightful outdoor meal or a game of cards under the stars. Elevate your camping experience with our reliable products, where quality meets the great outdoors!"

Task Oriented Prompts

- Focusing on a specific task helps the model to develop the output while taking that perspective into consideration



What is in this image?

The image shows a silver sedan that has been in an accident. The front end of the car is damaged, with the hood crumpled, and the bumper partially detached. The car is resting against a utility pole on the sidewalk, and there is visible smoke coming from the engine area, indicating a recent crash. The setting appears to be a residential street, as indicated by the presence of houses in the background, a neatly trimmed hedge, and a sidewalk.

Analyze this car damage video for an auto insurance report, focusing on identifying and detailing damage

In the image provided, we observe a silver sedan that has sustained significant front-end damage. The vehicle's hood is notably crumpled, and the front bumper has suffered substantial deformation, partially detaching from the chassis. The right headlamp assembly appears to be damaged or dislodged. There is also visible damage to the front grille area. The car is in contact with a utility pole, suggesting a frontal collision with the pole at a relatively low or moderate speed. Smoke is emanating from the engine compartment, suggesting potential mechanical or electrical damage that may need further inspection for a comprehensive assessment. The airbags do not appear to be deployed in this image. The vehicle will require a detailed inspection to evaluate the full extent of structural and mechanical damage for a complete insurance report.

Handling refusals

- Focusing on a specific task helps the model to develop the output while taking that perspective into consideration



Plan this meal

Sorry, I can't provide that information.

What information do you need ?

I apologize for the confusion earlier. To plan a meal based on the image provided, I would need to know the following:
The number of people dining to ensure there's enough food for everyone.

Any dietary restrictions or preferences to accommodate (e.g., vegetarian, gluten-free, allergies).

The occasion or type of meal being planned (e.g., casual lunch, dinner party).

Desired balance of the meal in terms of food groups and nutrition.

With this information, I could help suggest how to portion the dishes, create a balanced meal, and ensure the meal meets the needs of all diners.

Adding Examples

- Add examples that represent the type of responses you're looking for



Q: What is this dog breed and what is its average weight & height?

A: Breed: pomeranian; average weight: 3-7 lbs; average height: 8-14 inches

A: Breed: English bulldog; average weight: 50-55 lbs; average height: 14-16 inches



A: Breed: Border Collie; average weight: 30-45 lbs; average height: 18-22 inches

Breakdown Requests

Try breaking down complex requests step-by-step to create manageable sub-goals



You are an AI assistant that helps create shopping lists and instructions to prep food based on recipes. Please provide a list of ingredients to purchase, and then provide a list of instructions to prep the ingredients based on the recipe card.

Defining Output format

- Clearly mention the desired format for the output, such as markdown, JSON, HTML, etc.



You are an AI assistant that writes detailed listings for vacation rental properties. Summarize in three sentences the optimal features of the listing, based on only what is visible in the image. Conclude with one sentence on why the customer should book today.

Image Playground


- Generate compelling images from text and image prompts with models like gpt-image-1 and more.

← Images playground ▾

Help

Deployments gpt-image-1 ▾ Clear all outputs View code Filters feedback

List Grid



Describe the image you want to generate...



1024x1024



Quality

1v



Generate

Messages in the image playground are visible to anyone with access to this resource and using the API.

Video Playground



- Generate high-fidelity videos from latest models like Sora in Azure AI Foundry Models.

← Video playground PREVIEW

Create compelling video assets using text input. Select "Deploy now" to get started.

Pre-built prompts Generation history

Industry: All Deploy now



Lab – Prompt Engineering

Lab Description

- Explore various advanced Prompt Engineering Techniques